

FAST: framework for heterogeneous medical image computing and visualization

Erik Smistad · Mohammadmehdi Bozorgi · Frank Lindseth

Received: 12 December 2014 / Accepted: 30 January 2015
© CARS 2015

Abstract

Purpose Computer systems are becoming increasingly heterogeneous in the sense that they consist of different processors, such as multi-core CPUs and graphic processing units. As the amount of medical image data increases, it is crucial to exploit the computational power of these processors. However, this is currently difficult due to several factors, such as driver errors, processor differences, and the need for low-level memory handling. This paper presents a novel FrAme-work for heterogeneous medical image computing and visualization (FAST). The framework aims to make it easier to simultaneously process and visualize medical images efficiently on heterogeneous systems.

Methods FAST uses common image processing programming paradigms and hides the details of memory handling from the user, while enabling the use of all processors and cores on a system. The framework is open-source, cross-platform and available online.

Results Code examples and performance measurements are presented to show the simplicity and efficiency of FAST. The results are compared to the insight toolkit (ITK) and the visualization toolkit (VTK) and show that the presented framework is faster with up to 20 times speedup on several common medical imaging algorithms.

Conclusions FAST enables efficient medical image computing and visualization on heterogeneous systems. Code examples and performance evaluations have demonstrated

that the toolkit is both easy to use and performs better than existing frameworks, such as ITK and VTK.

Keywords GPU · Parallel · Medical · Image · Computing · Visualization · Heterogeneous · OpenCL

Introduction

An increasing amount of medical image data is becoming available for any given patient today. Modern image analysis techniques make it possible to extract and visualize more information from the images. The race for using the increasing amount of image data more effectively is paramount for better diagnostics and therapy in the future. Still, concurrent medical image computing and visualization of both static and dynamic real-time data is computationally expensive. In the efforts toward improving computer-assisted radiology and surgery, this may entail that computational demanding research methods that have been assessed to be quantitatively better than existing methods (in terms of accuracy for example) cannot be used in a routine clinical setting due to time constraints.

Most modern computer systems are heterogeneous in the sense that they consist of several different processors, such as multi-core CPUs and graphic processing units (GPUs). These processors enable parallel processing, which can accelerate many medical image computing tasks significantly [7,25]. The programming of this hardware is, however, still difficult due to several factors. One factor is that the software needed to use the hardware, such as GPU drivers and compilers, may contain errors which are hard to debug. Also, the different manufacturers may have interpreted the standards differently. This forces programmers to do more debugging and testing.

E. Smistad (✉) · M. Bozorgi · F. Lindseth
Department of Computer and Information Science,
Norwegian University of Science and Technology,
Sem Saelandsvei 7-9, 7491 Trondheim, Norway
e-mail: smistad@idi.ntnu.no

E. Smistad · F. Lindseth
SINTEF Medical Technology, Trondheim, Norway

Since the programmer cannot change proprietary software such as GPU drivers, the programmer may even have to write separate code for different hardware manufacturers and software versions. The result is increased software development overhead and fragmented source code.

GPUs were originally programmed using shaders intended for graphics rendering. Newer frameworks, such as CUDA [18], enable general-purpose programming of GPUs. The open computing language (OpenCL) [29] is an open standard for parallel programming of heterogeneous systems. OpenCL enables parallel programming of different processors such as multi-core CPUs and GPUs. These GPU programming tools expose the programmer to several hardware details. For instance, most GPUs have their own memory that is separate from the computer's main memory. This memory is often divided into several different memory spaces such as global, texture, and constant memory [19]. Thus, the programmer has to explicitly move data between the different memories during execution.

In this article, we propose a framework called FAST (FrAmework for heterogeneous medical image computing and visualization). This framework aims to make it easier to do efficient processing and visualization of medical images on heterogeneous systems. The framework is open-source and available online.¹ FAST is also cross-platform, supporting Windows, Mac OS X, and Linux. The authors believe that in order to achieve satisfactory performance in the more computational demanding medical applications, the framework has to cover the entire pipeline from reading and streaming data to visualizing the result on the screen. Thus, the framework currently includes methods for:

- Reading, writing and streaming image data in different formats.
- Image processing algorithms such as filtering, segmentation, and registration.
- Surface mesh extraction and rendering.
- Multi-volume and slice rendering.

The framework aims to be easy to use by utilizing common programming paradigms from popular toolkits, such as the insight toolkit (ITK) and the visualization toolkit (VTK), and hiding the details of memory handling from the user. Also, the framework has many tests and benchmarks which enable the user to make sure that all the hardware and software are working properly and gives the performance and accuracy necessary for a whole range of medical image processing applications. We acknowledge that there exist many medical image computing algorithms created using ITK and VTK. The framework therefore supports interoperability with these

frameworks, such that image data can be shared and pipelines from FAST, ITK, and VTK can be linked. This may ease the integration of FAST into existing applications.

Related work

ITK [9,10] and VTK [12,21] are two of the most commonly used frameworks for medical image analysis and visualization. ITK contains several image processing algorithms used in the medical domain, while VTK is mostly used for visualization. Several of the image processing filters in ITK and VTK support multi-threading for execution on multi-core CPUs. In this multi-threading model, the input image is split among a set of threads. Each subimage is processed individually, and the result is stitched together. These frameworks were not initially created with support for GPU acceleration, except GPU-based rendering. However, extensions have been proposed to enable such support [2,11]. The current version of ITK (4.6) includes GPU implementations of some algorithms such as thresholding, smoothing, and optical flow registration. However, these are implemented as separate modules which are only available if compiled with a specific flag.

The open computer vision library (OpenCV) [20] is another popular image processing and visualization framework. However, this framework focus primarily on 2D image processing and lack several features that are important in the medical imaging domain such as 3D image processing, medical image formats, and surface extraction. Still, OpenCV was designed for computational efficiency and with a strong focus on real-time applications. Several algorithms in OpenCV are implemented for the GPU using OpenCL.

While these frameworks provide accelerated processing more as an extension and as an optional feature, the FAST framework presented in this article has been designed with heterogeneous accelerated processing in mind from the start and it is part of the core of the framework. We believe this will result in a framework that is faster and easier to use.

MeVisLab [13,15] is a software which focus on rapid prototyping of medical image software using a visual programming interface. It also supports integration with ITK and VTK and has support for multi-threading. FAST, on the other hand, focuses on high-performance heterogeneous medical image computing and visualization and has currently no visual programming interface.

One framework that aims to aid the development of image processing algorithms for different GPUs is the heterogeneous image processing acceleration framework (HIPAcc) [14]. However, HIPAcc focus on the design of image processing algorithms and does not include visualization and registration.

¹ <http://github.com/smistad/FAST/>.

Outline

The next section describes the details of the framework. The result section presents code examples and performance benchmarks of common medical image computing pipelines on different systems. Finally, a discussion and conclusion is presented.

Methodology

The FAST framework consists of five main layers, as illustrated in Fig. 1. The bottom layer is the actual hardware, i.e., the CPUs and GPUs. The second layer are the drivers for this hardware, which are provided by the hardware manufacturers. Next is the library layer, which consists of several libraries that are needed in the framework. The libraries in this layer are:

- **Open Computing Library (OpenCL)**—An open standard for parallel programming on heterogeneous systems, including multi-core CPUs, GPUs, and FPGAs. It is supported by most processor manufacturers including AMD, NVIDIA, and Intel.
- **Open Graphics Library (OpenGL)**—A cross-platform library for visualization.
- **GL Extension Wrangler (GLEW)**—A library for handling OpenGL extensions.
- **Eigen**—A fast cross-platform linear algebra library.
- **Qt**—A cross-platform graphical user interface (GUI) toolkit.
- **Boost**—A C++ utility library.

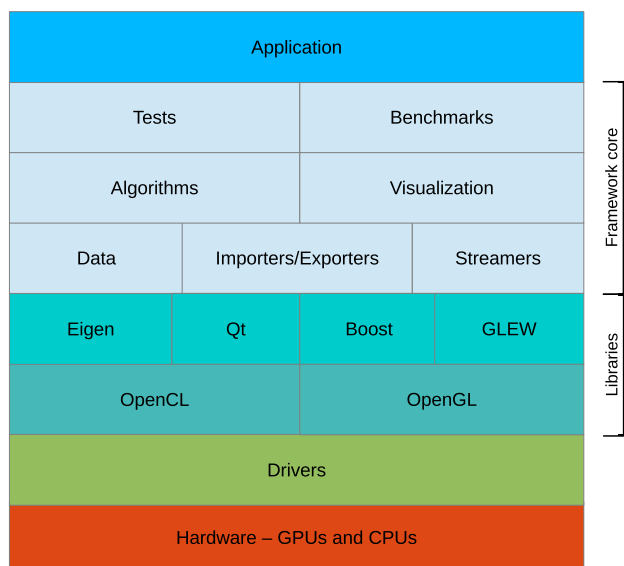


Fig. 1 Block diagram of the framework. The numbers indicate which section describes the different parts of the framework

The next layer is the core of the framework, which is split into several groups:

- **Data**—Objects for data (both static and dynamic) such as images and meshes, which enables the synchronized processing of such data on a set of heterogeneous devices.
- **Importers/Exporters**—Data import and export objects for different formats such as MetaImage (.mhd), raw, ITK, and VTK.
- **Streamers**—Objects that enable streaming of data.
- **Algorithms**—A set of commonly used filtering, segmentation, and registration algorithms.
- **Visualization**—A set of renderers such as image, volume, slice, and mesh renderers.
- **Tests**—A set of tests for the framework which ensures that all parts of the framework are working properly.
- **Benchmarks**—Mechanisms for measuring, assimilating, and reporting the performance of all operations in the framework.

The last layer is the application layer. The framework may be both a stand-alone application, which enables benchmarking and tests of a heterogeneous system, and an external library for other medical image computing applications.

The rest of this section will describe each part of the framework in more detail, but first the execution pipeline of the framework is described.

The execution pipeline

FAST uses a demand-driven execution pipeline similar to what is used in ITK and VTK. This entails that each processing step is first linked together to form a pipeline that is not executed until some object calls the update method. This can be done in two ways:

- Explicitly by calling the update method on an object in the pipeline.
- Implicitly by a renderer which calls update on its input connections several times per second.

The pipeline consists of process objects, which extend the abstract base class *ProcessObject*. A process object is an object that performs processing and may have zero, one, or several parent process objects. Most process objects produce data objects which extend the abstract base class *DataObject*. Similar to the newest version of VTK (version 6), FAST uses a pipeline where the data objects are not explicitly part of the pipeline. Figure 2 illustrates a simple pipeline with these two types of objects and how they are connected.

Data objects have an internal timestamp. The timestamp is always updated when the data are changed. Each process




Fig. 2 A simple pipeline with process (blue/bright) and data objects (orange/dark). The arrows indicate how the objects are connected

object has a list of timestamps for each connection. These timestamps represent which version of the data objects was used the last time the process object was executed. In addition, each process object has a flag indicating whether it has been modified or not. This could be a parameter or input change.

When the update method is called on a process object, it will first call update on all its parent objects. Thus, update will be called on all objects backwards in the pipeline until a process object with no input connections is encountered (e.g., an importer object). A process object will re-execute by calling its execute method, if it is modified or one of its input connections has changed timestamps. Thus, each process object will implement its own execute method, while the update method is the same for each process object.

Data management

Throughout this article, we will use the OpenCL terminology and refer to a processor with its memory as a *device* and the main CPU as the *host*.

Data organization and synchronization is one of the key components in the proposed framework. Image data are represented by an object called *Image* which is used for both 2D and 3D image data. These image objects represent an image on all devices, and its data are guaranteed to be coherent on any devices after being altered. Thus, if an image is changed on one device, it will also be changed on the other devices before the data are required on those devices. The same applies for other types of data like meshes where an object called *Mesh* is used to represent a mesh on all devices (see Fig. 3). Dynamic data, such as temporal 2D and 3D image data, are also supported. This is discussed in further detail in “Data streaming.”

Data access

Two forms of data access are possible in the framework: (1) read-only and (2) read and write. The general rule is that several devices can perform read operations on a data object




Fig. 3 A data object (e.g., an image) has pointers to all the devices where the data are stored. Using the OpenCL terminology, a device is a processor with its memory, and the host is the main CPU

at the same time. However, if a device needs to write data, only that device can have access to the data object at that time. This policy ensures data coherency across devices. Thus, if a device wants to write to an image, it has to wait for all other operations on that image to finish. When a device is writing to an image, no other devices can read or write to that image.

To enforce this policy, several *DataAccess* objects are introduced for each data object. For instance, in OpenCL, an image can be represented either as a buffer (i.e., a regular array) or as an image/texture. Thus, there exist one *OpenCLBufferAccess* object and one *OpenCLImageAccess* object to facilitate such access to image data. From these objects, an OpenCL *Image* or *Buffer* object can be retrieved, which is needed to perform OpenCL computations on the image. Access to the image from the main memory can also be requested for doing processing on the CPU using C++. The *DataAccess* objects also have methods for releasing the access, thus enabling other devices to perform write operations on the image. The access will also be released in the destructor of this object to avoid deadlocks. When the access is released, the OpenCL *Image/Buffer* object pointer is invalidated to ensure that the program can no longer manipulate the data. However, this does not delete the actual data on the device. When write access to an object is requested, the framework will check that any previous access objects have been released.

Data change

Every time data are changed on a device, the change should be reflected on the other devices as well. However, this does not have to be done immediately. Updating the data can be done the next time the data are requested on another device. This is often referred to as lazy loading. The benefit of lazy loading is that the number of data transfers can be reduced. However, the drawback is that there will be a transfer cost

the next time the data are requested on a device which does not have the updated data.

Thus, each data object has a set of flags indicating whether the data (in the form of OpenCL buffers, images and C++ pointers) are up to date for each device. When one device has changed some data, these are set to false for all devices except the device in which the change was performed. Next time the data are requested on a device, the flag is checked and if it is false, a data transfer will start and the flag will be set to true for that device.

Data removal

The amount of memory available on a system as well as on graphic cards is limited and may not be enough when working on large datasets. Thus, it is crucial to remove data that are not needed anymore. Data may be deleted explicitly by the programmer; however, this is a burden for the programmer and may easily be forgotten. After the entire pipeline has been defined by the programmer, it is known which process objects need which data objects as input. Thus, it is possible to delete a data object after all the process objects that use this data object have finished execution. This requires each process object to retain and release the data objects when they are defined as input and when the process object is finished using it. To facilitate this, each data object has a reference counter and when it reaches zero, the data are deleted.

Data types

Medical images are represented in different formats. Some common examples are: ultrasound (unsigned 8 bit integer), CT (signed/unsigned 16 bit integer), and MR (unsigned 16 bit integer). The framework currently supports the following data formats for images:

- TYPE_FLOAT—32 bit floating point number
- TYPE_UINT8—8 bit unsigned integer
- TYPE_INT8—8 bit signed integer
- TYPE_UINT16—16 bit unsigned integer
- TYPE_INT16—16 bit signed integer

An image can also have multiple channels, or components, and currently 1–4 channels are supported.

Data import and export

Data can be imported to and exported from the framework in several different forms such as:

- MetaImage file (.mhd, .raw, and .zraw)
- Image file (.jpg, etc.,)

- ITK image object
- VTK image object
- VTK file (.vtk)

In the future, the framework will also support common data formats such as DICOM [16] and NIfTI [17].

Data streaming

Streamers are process objects that provide access to dynamic data. This can for instance be real-time images from an ultrasound probe or a series of images stored on disk. The output of streamer objects is a *DynamicData* object, which has a method for retrieving the current frame in the stream. The *DynamicData* objects can contain one of several types of data such as images or meshes. The streamers read data into the *DynamicData* object in a separate thread so that processing and data streaming can be performed concurrently. Streamers can use one of three different streaming modes:

- STREAMING_MODE_NEWEST_FRAME_ONLY
This will only keep the newest frame in the *DynamicData* object.
- STREAMING_MODE_PROCESS_ALL_FRAMES
This will keep all frames in the *DynamicData* object, but will remove the frame from the object after it has been processed.
- STREAMING_MODE_STORE_ALL_FRAMES
This will store all frames in the *DynamicData* object.

For the second of these streaming modes, it is also important to limit the size of the dynamic data buffer so that the streaming does not use up all memory. With this mode, it is therefore possible to set the maximum size of the dynamic data buffer. A producer–consumer model is used to synchronize the use of the data.

The data and process objects are designed so that it is easy to accept both static and dynamic data as input and output to an algorithm.

Algorithms

Algorithms are implemented in the framework as process objects and thus have to override the execute method. Currently, only a few filtering, segmentation, and registration algorithms have been implemented such as Gaussian smoothing, seeded region growing [1], thresholding, skeletonization [8], iterative closest point [3], and surface extraction (marching cubes) [23]. All algorithms support parallel processing on CPUs and GPUs. In the near future, we plan to implement and integrate several other algorithms such as level set segmentation, Kalman filter object tracking [28], gradient vector flow [22,27], and tube detection filters [24,26].

Visualization

Graphical user interface and rendering

Qt is used in FAST as the graphical user interface. Qt is cross-platform, supports multi-threading, direct rendering from OpenGL and event handling of keyboard and mouse input. A visualization window in the FAST framework can have multiple views, and each view can have multiple renderers. Windows are implemented using Qt's *QWidget* class, while the *View* extends the *QGLWidget*, which is a widget that may be rendered to by OpenGL. The FAST renderers do the actual rendering. These renderers and the event handling are executed in one thread, while the pipeline is run in another thread. This enables concurrent visualization, camera movement, and pipeline execution. Five different types of renderers are currently available in FAST:

- Image renderer—For displaying 2D images.
- Slice renderer—Extracts and displays an image from a volume in an arbitrary plane using trilinear interpolation.
- Mesh renderer—Renders a mesh.
- Volume renderer—Creates an image of a volume using ray casting [4].
- Point renderer—Renders a list of points.

Scene graph

Correct placement of images and geometry in the visualization scene is important. FAST uses a scene graph for this purpose. In this directed graph, each data object has a node. All data nodes are connected to a parent node, which can be another data node or a root node. Each edge between the nodes has a transformation object. This transformation determines how data are positioned relative to other nodes. Figure 4 shows an example of a scene graph with three




Fig. 4 An example of a scene graph. Images A and B are registered because they share a root node. Image C is not registered to any other data. Each edge between the nodes has a transformation object. This transformation determines how data are positioned relative to other nodes. Meshes A and B are dependent on Image A. Thus, moving Image A will also move these meshes

images and two meshes. The images A and B share a root node and are therefore registered. Image C is not registered to any other data. Image A is placed in the visualization scene by applying the transformation T_{IA} to the image. Similarly, image B uses the transformation T_{IB} . Since these images are registered, the corresponding voxel position in image B of a voxel position in image A can be determined by first applying the transformation T_{IA} , and then the inverse transformation of image B T_{IB}^{-1} . The meshes A and B are related to image A. These meshes may for instance be the result of a segmentation of image A. Mesh A is placed in the visualization scene by first applying the transformation from the mesh to image A T_{MA} and then the transformation from image A to the root node T_{IA} . Thus, if image A is moved in the scene, the meshes A and B are also moved.

When an image or mesh is created, a corresponding data node is created in the scene graph and connected to a root node. However, if the data are created from another data object, it is connected to the data node of that data object instead. For instance, the surface extraction algorithm will connect the resulting mesh to the image used to create the mesh. A visualization of an image and a segmented surface mesh using the scene graph are shown in Fig. 6. When importing a MetaImage, any transformation information such as translation and rotation is read from the MetaImage file (.mhd) and put in the scene graph.

Tests

As much as possible of the framework should be covered by unit and system tests. This enables a user to ensure that the framework is working correctly on the user's current software and hardware configuration. The authors know by experience that new drivers, compilers, and libraries can introduce errors that may stop the framework from working properly. These tests enable a user to quickly detect these problems. The tests are written using the Catch C++ testing framework [5]. Realistic test data are needed to test the framework properly and are therefore provided for download.²

The framework is available on the open-source community Web site GitHub. Each time a user contributes to the project, three different computers will execute all tests with the new code and verify that everything is working. These machines use all the supported operating systems Windows, Mac OS X, and Ubuntu Linux and processors from Intel, AMD, and NVIDIA. Thus, the source code of the framework is tested continuously on several hardware and software configurations. We believe this is needed in order to ensure the stability of FAST.

² <http://github.com/smistad/FAST/wiki/Test-data>.

Benchmarks

Users may also want to test how well their current setup performs and see how performance changes when software and hardware changes are introduced. For this purpose benchmarks are provided, which are tests of different pipelines in which performance is measured and reported.

Results

This section first presents some examples of how the framework can be used. These examples are provided to show how easy it is to set up pipelines in FAST. Next, the performance of the framework is measured and compared to that of ITK and VTK.

Code examples

The first example is a simple pipeline of four steps: import 3D image from disk, Gaussian smoothing, surface extraction, and rendering. The result is shown in Fig. 5. The steps of the pipeline are linked together using the *getOutputPort* and *setInputConnection* methods of the process objects. This is the same method used by ITK and VTK. The *::pointer* types are smart pointers which are created with the *New* method. These pointers reduce memory problems such as memory leakage.

Example 1: Pipeline A

```
// Import image
ImageFileImporter::pointer importer =
    ImageFileImporter::New();
importer->setFilename("image.mhd");

// Blur image with Gaussian smoothing
GaussianSmoothing::pointer smoothing =
    GaussianSmoothing::New();
smoothing->setInputConnection(
    importer->getOutputPort());
smoothing->setStandardDeviation(1.0);

// Extract surface mesh with marching cubes
SurfaceExtraction::pointer extraction =
    SurfaceExtraction::New();
extraction->setInputConnection(
    smoothing->getOutputPort());

// Render surface mesh
MeshRenderer::pointer meshRenderer =
    MeshRenderer::New();
meshRenderer->addInputConnection(
    extraction->getOutputPort());

// Render slice
SliceRenderer::pointer sliceRenderer =
    SliceRenderer::New();
sliceRenderer->addInputConnection(
    smoothing->getOutputPort());
sliceRenderer->setSlicePlane(PLANE_X);

// Create a window, attach the renderers and
// start pipeline
SimpleWindow::pointer window =
    SimpleWindow::New();
window->addRenderer(meshRenderer);
window->addRenderer(sliceRenderer);
window->start();
```




Fig. 5 Result of pipeline A in Example 1. A 3D ultrasound image is first smoothed. Then, surface extraction is used to extract a surface mesh from the smoothed image. Finally, a slice of the smoothed 3D image is rendered together with the surface mesh

This pipeline can easily be changed from using a single static image as input to a stream of images by only substituting the *Importer* object with a *Streamer* object. The rest of the pipeline is the same. Example 2 shows how a *ImageFileStreamer* object is created to stream a series of MetaImages from disk. The streamer object uses a filename format to find files. The hash sign (#) is replaced by an integer index which changes for each image that is loaded. It is possible to change the start index and step which are 0 and 1, respectively, by default. The streamer stops when no more images with the format are found.

Example 2: Streaming images

```
ImageFileStreamer::pointer streamer =
    ImageFileStreamer::New();
streamer->setFilenameFormat("image_frame_#.mhd");
```

The user may want to specify which device should be used as the default device. This is done using the *DeviceManager* object as shown in Example 3. However, each process object may override this if desired.

Example 3: Set the default device to be a GPU

```
DeviceManager::setDefaultDevice(
    DeviceManager::getOneGPUDevice());
```

The next example shows another pipeline. This pipeline performs region-growing segmentation on an image, extracts the surface mesh of the segmentation, and finally renders the mesh and a slice of the input image. The result can be seen in Fig. 6.

Example 4: Pipeline B

```
// Import image
ImageFileImporter::pointer importer =
    ImageFileImporter::New();
importer->setFilename("CT-Abdomen.mhd");
```




Fig. 6 Result of pipeline B in Example 4. Region growing is used to segment the bone structure from a CT scan. A surface mesh is extracted from the segmentation and rendered together with a slice of the CT scan. The scene graph is used to correctly position the two data objects.

```
// Segment image with region growing
SeededRegionGrowing::pointer segmentation =
    SeededRegionGrowing::New();
segmentation->setInputConnection(
    importer->getOutputPort());
segmentation->addSeedPoint(261,284,208);
segmentation->setIntensityRange(150, 5000);

// Extract surface mesh with marching cubes
SurfaceExtraction::pointer extraction =
    SurfaceExtraction::New();
extraction->setInputConnection(
    segmentation->getOutputPort());

// Render slice plane
SliceRenderer::pointer sliceRenderer =
    SliceRenderer::New();
sliceRenderer->setPlaneToRender(PLANE_Z);
sliceRenderer->setIntensityWindow(1000);
sliceRenderer->setIntensityLevel(0);
sliceRenderer->setInputConnection(
    importer->getOutputPort());

// Render surface mesh
MeshRenderer::pointer meshRenderer =
    MeshRenderer::New();
meshRenderer->addInputConnection(
    extraction->getOutputPort());

// Create a window, attach the renderers and
// start pipeline
SimpleWindow::pointer window =
    SimpleWindow::New();
window->addRenderer(sliceRenderer);
window->addRenderer(meshRenderer);
window->start();
```

Pipeline C imports a 2D image and performs binary threshold segmentation. The segmentation is skeletonized and finally rendered using the ImageRenderer as shown in Fig. 7.

Example 5: Pipeline C

```
// Import image
ImageFileImporter::pointer importer =
    ImageFileImporter::New();
```




Fig. 7 Result of pipeline C in Example 5 where an image of the retina blood vessels is thresholded and skeletonized using iterative thinning

```
importer->setFilename("image.png");

// Segment image with thresholding
BinaryThresholding::pointer thresholding =
    BinaryThresholding::New();
thresholding->setInputConnection(
    importer->getOutputPort());
thresholding->setLowerThreshold(0.5);

// Skeletonize the segmentation
Skeletonization::pointer skeletonization =
    Skeletonization::New();
skeletonization->setInputConnection(
    thresholding->getOutputPort());

// Render image
ImageRenderer::pointer renderer =
    ImageRenderer::New();
renderer->addInputConnection(
    skeletonization->getOutputPort());

// Create a window, attach the renderers and
// start pipeline
SimpleWindow::pointer window =
    SimpleWindow::New();
window->addRenderer(renderer);
window->start();
```

The next pipeline first imports two point sets from VTK files (.vtk). The *PointSet* object is a data object, which only contains a set of points. These point sets are then registered using the iterative closest point (ICP) Algorithm [3]. Finally, the point sets are rendered using the *PointRenderer* (see Fig. 8).

Example 6: Pipeline D

```
// Import two point sets
PointSetImporter::pointer importerA =
    PointSetImporter::New();
importerA->setFilename("pointsA.vtk");
PointSet::pointer pointsA =
    importerA->getOutputPort();
PointSetImporter::pointer importerB =
    PointSetImporter::New();
```





Fig. 8 Two point sets of pipeline D in Example 6 before and after the iterative closest point (ICP) algorithm are used to register the two sets

```
importerB->setFilename("pointsB.vtk");
PointSet::pointer pointsB =
    importerB->getOutputPort();

// Run iterative closest point
IterativeClosestPoint::pointer icp =
    IterativeClosestPoint::New();
icp->setMovingSet(pointsA);
icp->setFixedSet(pointsB);

// Render the two point sets
PointRenderer::pointer renderer =
    PointRenderer::New();
renderer->addInput(pointsA, Color::Blue(), 10);
renderer->addInput(pointsB, Color::Green(), 5);

// Create a window, attach the renderers and
// start pipeline
SimpleWindow::pointer window =
    SimpleWindow::New();
window->addRenderer(renderer);
window->start();
```

Performance

The runtime and memory usage of pipeline A, B, C, and D (see Examples 1, 4, 5 and 6) were measured and collected in Table 1. The runtime is the average of 10 runs on each system. The memory usage is measured using the system monitor and includes only the system memory and not the GPU memory usage. The same pipelines were implemented and measured in ITK and VTK for comparison. Several of the ITK and VTK image processing filters used in these pipelines including smoothing, thresholding, thinning, region growing, and iterative closest point use multi-threading. Three different computer systems, all with solid-state drives (SSD), were used for the measurements:

- Intel i5 3.4GHz CPU with 16GB RAM, NVIDIA GeForce GTX 970 4 GB running Windows 8.1.
- AMD A10 CPU with 16GB RAM, AMD Radeon R9 290 GPU 4GB running Ubuntu 14.04 Linux.
- Intel i5 3.4GHz CPU with 16GB RAM, NVIDIA GeForce GTX 780M 4 GB running Mac OS X 10.9.

The following datasets were used for the different pipelines, and informed consent was obtained from all patients for being included in the study:

- **Pipeline A:** 3D ultrasound image, unsigned 8 bit integer, $276 \times 249 \times 200$ voxels ≈ 14 MB.
- **Pipeline B:** CT image, signed 16 bit integer, $512 \times 512 \times 426$ voxels ≈ 223 MB.
- **Pipeline C:** 2D image, 565×584 pixels ≈ 11 kB.
- **Pipeline D:** Two point set files of the left ventricle with 386 3D points each of ≈ 31 kB.

Discussion

As more medical imaging data becomes available, a framework that exploits the increasingly heterogeneous and parallel computers is needed. These heterogeneous systems are hard to program due to several factors such as such as driver errors, processor differences, and the need for low-level memory handling. The FAST framework makes medical image processing and visualization easier by using familiar programming paradigms and hiding the details of memory handling from the user, while still enabling the use of all processors and cores on a system. Errors and differences in proprietary software and hardware (e.g., GPU drivers) cannot be fixed by the medical imaging community, as the development of these is dependent on the manufacturers. FAST aims to provide a large set of tests and benchmarks to detect and report these problems. This enables an easy way for a user to check if there are any problems for a specific setup.

Although ITK has a couple of algorithms that support OpenCL as an optional extension, we believe that it is necessary to support heterogeneous processing in the entire framework to achieve the best performance. This include all steps in a pipeline from data import, to processing and visualization. Enabling this kind of support in ITK, VTK, or MeVisLab would most likely mean rewriting the entire core of these toolkits.

The code examples in the previous section show how easy it is to set up different pipelines consisting of data import, streaming, processing, and rendering. Implementation of the same pipelines in ITK and VTK required more lines of code and code complexity mainly due to the need for exporting data from ITK to VTK and templating in ITK. One may argue that this is because ITK and VTK have more features.

Table 1 Performance measurements of the four pipelines in Examples 1, 4, 5 and 6

Pipeline	Framework	System (CPU/GPU/OS)	Runtime (ms)	Memory usage
Pipeline A Data import, Gaussian smoothing, surface extraction and rendering	FAST	Intel/NVIDIA/Windows	86 (0.2, 34, 52, 0.6)	92 MB
		AMD/AMD/Linux	52 (0.2, 23, 29, 0.6)	53 MB
	ITK and VTK	Intel/NVIDIA/Mac	135 (0.1, 71, 63, 0.5)	34 MB
		Intel/NVIDIA/Windows	696 (9, 97, 511, 78)	260 MB
		AMD/AMD/Linux	1133 (6, 262, 568, 295)	255 MB
		Intel/NVIDIA/Mac	704 (29, 68, 441, 165)	140 MB
Pipeline B Data import, region growing, surface extraction and rendering	FAST	Intel/NVIDIA/Windows	2293 (230, 1954, 108, 0.7)	398 MB
		AMD/AMD/Linux	2270 (262, 1848, 158, 0.8)	402 MB
	ITK and VTK	Intel/NVIDIA/Mac	5103 (273, 4588, 242, 0.7)	355 MB
		Intel/NVIDIA/Windows	3041 (346, 732, 1673, 290)	1.5 GB
		AMD/AMD/Linux	4615 (301, 906, 2309, 1099)	1.4 GB
		Intel/NVIDIA/Mac	3473 (765, 623, 1616, 467)	1.1 GB
Pipeline C Data import, thresholding, skeletonization and rendering	FAST	Intel/NVIDIA/Windows	50 (39, 4, 3, 3)	79 MB
		AMD/AMD/Linux	20 (9, 2, 6, 3)	55 MB
	ITK and VTK	Intel/NVIDIA/Mac	26 (11, 3, 9, 3)	42 MB
		Intel/NVIDIA/Windows	856 (3, 1, 819, 33)	28 MB
		AMD/AMD/Linux	489 (5, 0.8, 384, 99)	23 MB
		Intel/NVIDIA/Mac	721 (10, 0.4, 682, 28)	50 MB
Pipeline D Data import, iterative closest point and rendering	FAST	Intel/NVIDIA/Windows	25 (2, 23, 0.2)	80 MB
		AMD/AMD/Linux	21 (4, 17, 0.2)	52 MB
	ITK and VTK	Intel/NVIDIA/Mac	9 (1, 8, 0.2)	16 MB
		Intel/NVIDIA/Windows	293 (31, 84, 178)	22 MB
		AMD/AMD/Linux	241 (4, 109, 128)	20 MB
		Intel/NVIDIA/Mac	152 (6, 61, 85)	14 MB

The same pipelines were implemented in ITK and VTK for comparison. Three systems with different operating system and hardware were used. The runtime of each step in the pipeline is listed in parentheses

However, we believe that common operations, such as these four pipelines, should require little code.

The performance measurements in Table 1 and Figs. 9 and 10 show that FAST is faster for the four pipelines with speedups of up to 20 times compared to ITK and VTK. This speedup is mainly due to the fact that FAST is able to use the GPU for processing and rendering, while ITK and VTK rely on multi-threading for acceleration. All steps in the pipeline, including data import, processing, and rendering, are faster with the proposed framework, with region growing, 2D thresholding and data import in pipeline C as the only exceptions. As shown in Fig. 10, the largest difference in runtime is with rendering, where FAST uses from 0.2 to 3 ms, while VTK uses 28–1099 ms.

ITK and VTK use more system memory than FAST on pipelines A and B which use large 3D images as input with sizes of 14 and 223 MB. However, FAST uses the GPU for most of this processing and will therefore also use more GPU memory. The overall goal is to significantly reduce the total

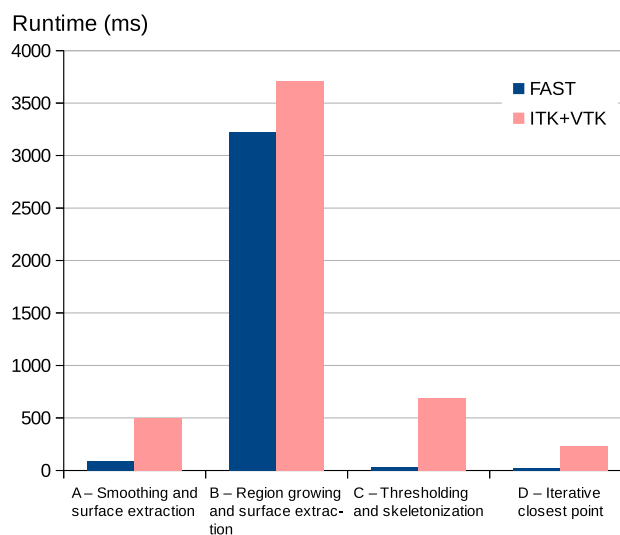


Fig. 9 Average runtime performance for all computer system in section “Performance” of the four pipelines in Examples 1, 4, 5, and 6 using FAST, ITK, and VTK.




Fig. 10 Detailed average runtime performance for all computer system in section “Performance” of the four pipelines in Examples 1, 4, 5, and 6 using FAST, ITK, and VTK.

memory usage in medical image computing and visualization applications. This will be done by avoiding data duplications and keeping tracking of the CPU and GPU memory used in future versions of FAST.

For FAST to be accepted by the medical imaging community, it needs commonly used algorithms to speed up the development of high-level image processing algorithms. We plan to implement more algorithms in FAST and hope that others will contribute to this open-source framework through the collaboration platform GitHub. Still, developers can use algorithms that already exist in ITK, as FAST supports interoperability with ITK pipelines. Documentation and examples are also vital for the framework’s success. Thus, an online open-source wiki has been created.³

The most important application for high-performance medical image processing and visualization is image guided surgery, where preoperative data are combined with intra-operative data which needs to be acquired, processed and visualized in the operating room. Optical and electromagnetic tracking are important in this context. Our future work will be on incorporating tracking in the framework to enable surgical navigation.

The authors strongly believe in open-source code for medical image computing and visualization as a mean for advancing the state of the art, as well as open data for evaluating algorithms. Each year a vast amount of new methods and modifications of existing ones are proposed (e.g., registration and segmentation algorithms). In order to increase the amount of methods that are actually used clinically, it is crucial that new algorithms are benchmarked thoroughly, both in terms of accuracy and computational performance. In the last decade, several challenges have been arranged [6], and open databases with an established ground truth are probably the best tool we have today to achieve this. As the proposed framework becomes more mature, the authors hope

that FAST can contribute to this effort and make more algorithms clinically ready faster.

Conclusion

A novel framework for efficient medical image computing and visualization has been presented. The framework was built from ground up with optimal performance on heterogeneous systems in mind. Code examples and performance evaluations have demonstrated that the toolkit is both easy to use and performs better than existing frameworks, such as ITK and VTK. Built-in benchmarking support will make additional fine tuning a lot easier and produce new insight about heterogeneous computing in the medical domain. As more quality and performance benchmarked functionality is added to the framework, the authors hope that FAST will be a valid tool for bringing more medical imaging software into clinical practice in the years to come.

Acknowledgments This project has received funding from the European Union’s Seventh Framework Programme for research, technological development and demonstration under Grant Agreement No. 610425. The hardware used in this project was funded by the MedIm (Norwegian Research School in Medical Imaging) Travel and Research Grant.

Conflict of interest Erik Smistad, Mohammadmehdi Bozorgi and Frank Lindseth declare that they have no conflict of interest.

References

- Adams R, Bischof L (1994) Seeded region growing. *IEEE Trans Pattern Anal Mach Intell* 16(6):641–647
- Beare R, Micevski D, Share C, Parkinson L, Ward P, Goscinski W, Kuiper M (2011) CITK - an architecture and examples of CUDA enabled ITK filters. *Insight J* 2011(Jan-June):1–8
- Besl PJ, McKay ND (1992) A method for registration of 3-D shapes. *IEEE Trans Pattern Anal Mach Intell* 14(2):239–256

³ <http://github.com/smistad/FAST/wiki/>.

4. Bozorgi M, Lindseth F (2014) GPU-based multi-volume ray casting within VTK for medical applications. *Int J Comput Assist Radiol Surg*. doi:10.1007/s11548-014-1069-x
5. Catch. C++ automated test cases in headers. <https://github.com/philsquared/Catch/>. Accessed 10 Oct 2014
6. Consortium for open medical image computing. Grand challenges in biomedical image analysis. <http://grand-challenge.org/>. Accessed 25 Nov 2014
7. Eklund A, Dufort P, Forsberg D, Laconte SM (2013) Medical image processing on the GPU—Past, present and future. *Med Image Anal* 17(8):1073–1094
8. Gonzalez RC, Woods RE (2008) Digital image processing, 3rd edn. Pearson Prentice Hall, New Jersey
9. Ibanez L, Schroeder W (2004) The ITK software guide, 2.4th edn. Kitware, Chapel Hill
10. Kitware. Insight toolkit (ITK). <http://itk.org/>. Accessed 18 Aug 2014
11. Kitware. ITK release 4 GPU acceleration. http://www.itk.org/Wiki/ITK/Release_4/GPU_Acceleration/. Accessed 10 Oct 2014
12. Kitware. Visualization toolkit (VTK). <http://www.vtk.org/>. Accessed 18 Aug 2014
13. Koenig M, Spindler W, Rexilius J, Jomier J, Link F, Peitgen H-O (2006) Embedding VTK and ITK into a visual programming and rapid prototyping platform. In: Proceedings of SPIE, vol. 6141, pp 61412O–61412O-11
14. Membarth R, Hannig F, Teich J, Körner M, Eckert W (2012) Generating device-specific GPU code for local operators in medical imaging. In: Proceedings of the 26th IEEE international parallel & distributed processing symposium (IPDPS), number section III
15. MeVis Medical Solutions AG. MeVisLab. <http://www.mevislab.de>. Accessed 26 Jan 2015
16. Mildenerger P, Eichelberg M, Martin E (2002) Introduction to the DICOM standard. *Eur Radiol* 12:920–927
17. Neuroimaging Informatics Technology Initiative. NIFTI-1 data format. <http://nifti.nimh.nih.gov/>. Accessed 26 Jan 2015
18. NVIDIA Corporation. CUDA. <http://developer.nvidia.com/cuda-zone/>. Accessed 26 Jan 2015
19. Owens J, Houston M, Luebke D, Green S, Stone J, Phillips J (2008) GPU computing. In: Proceedings of the IEEE 96(5):879–899
20. Pulli K, Baksheev A, Korniyakov K, Eruhimov V (2012) Real-time computer vision with OpenCV. *Commun ACM* 55(6):61
21. Schroeder W, Martin K, Lorensen B (2006) Visualization toolkit: an object-oriented approach to 3D graphics, 4th edn. Kitware, Chapel Hill
22. Smistad E, Elster AC, Lindseth F (2012) Real-time gradient vector flow on GPUs using OpenCL. *J Real-Time Image Process*, 1–8
23. Smistad E, Elster AC, Lindseth F (2012) Real-Time Surface Extraction and Visualization of Medical Images using OpenCL and GPUs. *Norsk informatikkonferanse*, 141–152. Akademika forlag
24. Smistad E, Elster AC, Lindseth F (2014) GPU accelerated segmentation and centerline extraction of tubular structures from medical images. *Int J Comput Assist Radiol Surg* 9(4):561–575
25. Smistad E, Falch TL, Bozorgi M, Elster AC, Lindseth F (2015) Medical image segmentation on GPUs—a comprehensive review. *Med Image Anal* 20(1):1–18
26. Smistad E, Lindseth F (2014) A new tube detection filter for abdominal aortic aneurysms. In: Proceedings of MICCAI 2014 workshop on abdominal imaging: computational and clinical applications
27. Smistad E, Lindseth F (2014) Multigrid gradient vector flow computation on the GPU. *J Real-Time Image Process*
28. Smistad E, Lindseth F (2014) Real-time tracking of the left ventricle in 3D ultrasound using kalman filter and mean value coordinates. In: Proceedings MICCAI challenge on echocardiographic three-dimensional ultrasound segmentation (CETUS), pp 65–72, Boston
29. The Khronos Group. OpenCL. <http://www.khronos.org/opencl/>. Accessed 26 Jan 2015